

# Abstraction in Decision-Makers with Limited Information Processing Capabilities

---

Tim Genewein<sup>1,2,3</sup>, Daniel A. Braun<sup>1,2</sup>

<sup>1</sup>Max Planck Institute for Intelligent Systems

<sup>2</sup>Max Planck Institute for Biological Cybernetics

<sup>3</sup>Graduate Training Center for Neuroscience  
Tübingen, Germany



# Abstraction

- Neglecting **irrelevant** information with respect to a particular **purpose**
  - Example: chair for the purpose of sitting down
  - separation of structure from noise
  - treating different things as equal
- Crucial if information processing capabilities are limited
- Traditionally, the formation of abstractions is associated with high computational cost (e.g. clustering)
- We argue that abstractions arise as a consequence of limited computational capacity
- ... which could be induced by time- or memory-constraints, limited sample size or high signal-to-noise ratio

# Multi-Task Decision-Making scenario

- Here, we consider the following multi-task problem - the qualitative findings carry over to more general inference or decision-making scenarios
- Each task is formalized via a (relative) utility function  $U$
- Agent emits an action  $x$  conditioned on an observation  $y$ :  $p(x|y)$
- Assumptions:
  - Action-space is shared among all tasks
  - Observations are **fully** informative about the current task
  - Discrete, i.i.d. setting
- Multi-task abstractions – different tasks will be treated equally

# Thermodynamic Model for DM

- Decision-making with information processing limits
- Behavior/policy is modeled by a probability distribution
- Trade-off between gain in expected utility and transformation cost (from prior behavior  $p_0$  to optimized policy  $p(x|y)$ )

$$\operatorname{argmax}_{p(x|y)} \mathbf{E}_{p(x|y)}[U_y(x)] - \frac{1}{\beta} D_{\text{KL}}(p(x|y) || p_0(x))$$

- Variational problem has very similar mathematical form as a *free-energy difference* minimization
- Closed-form solution:

$$p(x|y) = \frac{1}{Z} p_0(x) e^{\beta U_y(x)}$$

# Temperature as rationality-parameter

$$\operatorname{argmax}_{p(x|y)} \mathbf{E}_{p(x|y)}[U_y(x)] - \frac{1}{\beta} D_{\text{KL}}(p(x|y) || p_0(x))$$

$$p(x|y) = \frac{1}{Z} p_0(x) e^{\beta U_y(x)}$$

- Limits:
  - Fully rational actor:  $\beta \rightarrow \infty$
  - Fully bounded actor:  $\beta \rightarrow 0$
- Normative framework for changing from **prior** belief/behavior to **posterior** belief/behavior with information processing cost
  - Bayes rule can be recovered as a special case

# Rate-Distortion for Decision-Making

- Extend free energy model by taking the average over observations and optimizing over the prior as well:

$$\operatorname{argmax}_{p_0(x)} \sum_y p(y) \left[ \operatorname{argmax}_{p(x|y)} \mathbf{E}_{p(x|y)} [U_y(x)] - \frac{1}{\beta} D_{\text{KL}}(p(x|y) || p_0(x)) \right]$$

- ... which can be rewritten:

$$\operatorname{argmax}_{p(x|y)} \underbrace{\sum_{x,y} p(x,y) U(x,y)}_{\mathbf{E}_{p(x,y)} [U]} - \frac{1}{\beta} \underbrace{\sum_y p(y) D_{\text{KL}}(p(x|y) || p(x))}_{I(x;y)}$$

- Same mathematical form as rate-distortion problem

# Rate-Distortion for Decision-Making

$$\operatorname{argmax}_{p(x|y)} \underbrace{\sum_{x,y} p(x,y) U(x,y)}_{\mathbf{E}_{p(x,y)}[U]} - \underbrace{\frac{1}{\beta} \sum_y p(y) D_{\text{KL}}(p(x|y) || p(x))}_{I(x;y)}$$

- Trade off between **high expected utility** and **low mutual information** between actions and observations (tasks)
- Optimal prior implies minimal average transformation cost for adapting to a particular task
  - If memory is limited,  $p(x)$  would be the optimal information to store
- Rate-distortion – a framework for lossy compression
  - Lossy compression requires neglecting “the right” information
  - Separation of structure and noise

# Temperature as rationality-parameter

- Well known (self-consistent) solution:

$$p(x|y) = \frac{1}{Z} p(x) e^{\beta U_y(x)}$$
$$p(x) = \sum_y p(y) p(x|y)$$

- Temperature governs trade-off between high expected utility and low task-specificity.
  - $\beta \rightarrow \infty$ : optimal adaptation to each task
  - $\beta \rightarrow 0$ : full abstraction – same policy for all tasks
- Agent with limited information processing capabilities (finite  $\beta$ ) can not adapt optimally to each task
  - Emergence of robustness/abstraction
  - Temperature governs the trade-off and thus the granularity of the abstraction



# Experiment I


- Four actions, two tasks
- Two temperature values
  - High inverse temperature: rational actor
  - Low inverse temperature: bounded actor

$x$			$\beta = 100$			$\beta = 1$		
			$p(x)$	$p(x y_1)$	$p(x y_2)$	$p(x)$	$p(x y_1)$	$p(x y_2)$
$[0, 0]$	$U(x, y_1)$	$U(x, y_2)$	0	0	0	0	0	0
$[0, 1]$	0	1	0.5	0	1	0	0	0
$[0.7, 0]$	0.7	0.7	0	0	0	1	1	1
$[1, 1]$	1	0	0.5	1	0	0	0	0

- Fully rational actor: selects maximum utility action in each task
- Fully bounded actor: always chooses suboptimal action
  - $\rightarrow$  mutual information between action and observation is zero

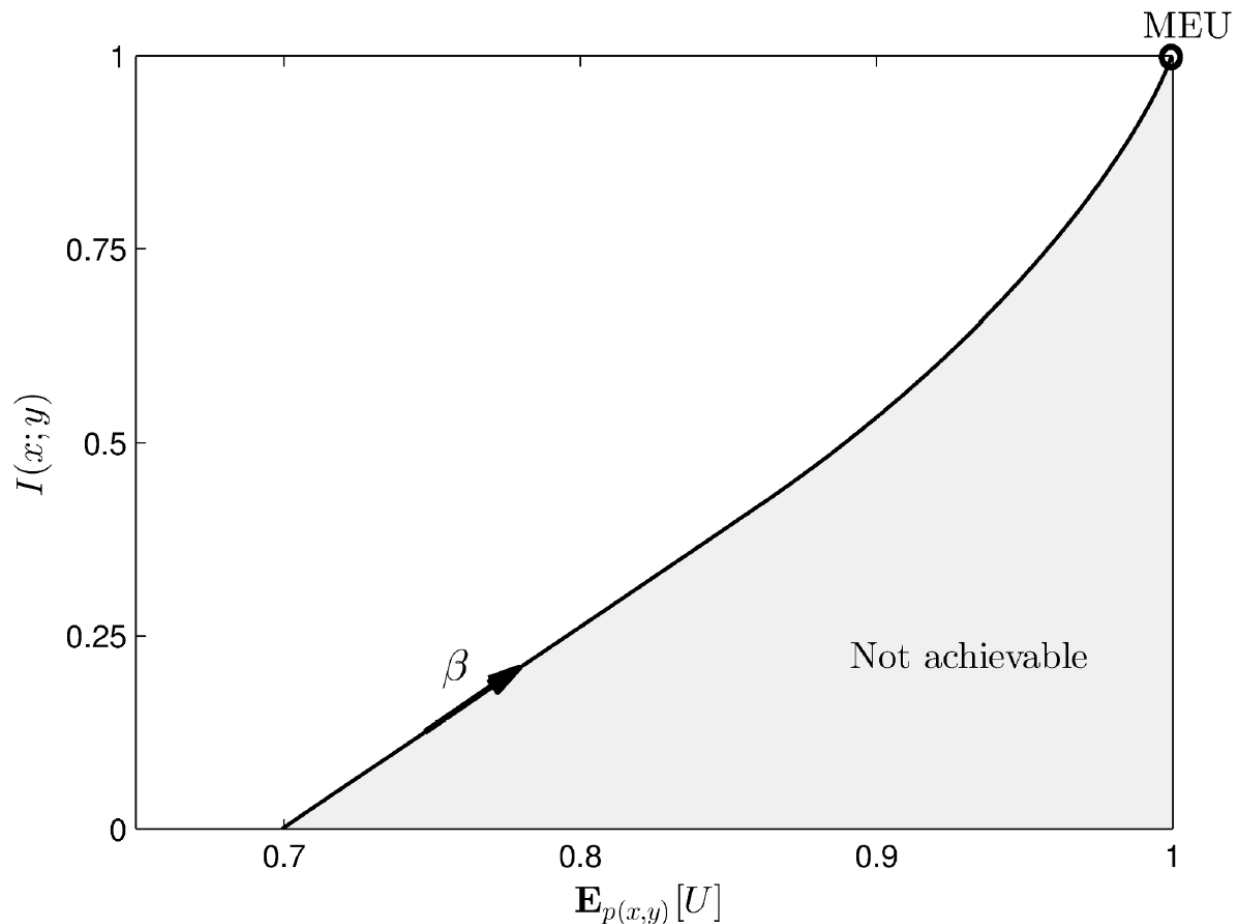
# Experiment I

- Initial mixing of maximum utility actions (increase in stochasticity)
- Gradual decrease of MI up to a “phase-transition”



# Experiment I

- Bounded-optimal actors lie on the rate-utility curve
- The shaded region is theoretically impossible to achieve



# Experiment II

- Actions are binary colorings of  $3 \times 3$  grids
- 512 possible actions/patterns
- Three environments
  - 1) One row and one column has to be all-white, each colored pixel yields +1 utility
  - 2) Any pattern with exactly four colored pixels yields +4 utility
  - 3) Any pattern with an even number of colored pixels yields +1 utility for each colored pixel
- Effective cardinality of the action-distribution is affected by boundedness of the actor


# Experiment II

$$\beta = 10$$

- Samples from conditionals and prior


A

Task 1




B

Task 2




C

Task 3



D

Prior




# Experiment II

$$\beta = 0.1$$

- Samples from conditionals and prior


A

Task 1




B

Task 2




C

Task 3



D

Prior



# Discussion

- Formalism also carries over to more general decision-making or inference cases
- Similar work
  - VAN DIJK, S. G. & POLANI, D. (2013). Informational Constraints-Driven Organization in Goal-Directed Behavior. *Advances in Complex Systems*.
  - STILL, S & CRUTCHFIELD, J. P. (2008). Structure or Noise? [arXiv:0708.0654v2](https://arxiv.org/abs/0708.0654v2) [physics.data-an]
  - Information Bottleneck Method, Relevant Information
  - Rational Inattention
- Certain regularizers might be elegant ways of implementing rate-distortion
- Continuous cases
  - No closed-form solutions for self-consistent equations
- Sampling representation
  - Representing prior and conditionals with samples
- Modelling **several layers** of abstraction
  - Unclear how to break down utility functions across layers